

Project_Py_ML_LR_Evaluation_Beer Consumption Analysis

The objective of this project will be to demonstrate the impacts of variables on beer consumption in a given region and the consumption forecast for certain scenarios

Note-Rename the column names Respectively:

Data: 'Date', Temperatura Media (C): 'Medium Temp', Temperatura Minima (C): 'Min Temp', Temperatura Maxima (C): 'Max Temp', Precipitacao (mm): 'Precipitation(mm)', Final de Semana: 'End of week', Consumo de cerveja (litros): 'Beer consumption (liters)'

importing the required libraries

```
In [230... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

load the given dataset

```
In [231... df = pd.read_csv("Project_Py_ML_LR_Evaluation.csv")
df.head()
```

```
Out[231... 
```

| | Data | Temperatura Media (C) | Temperatura Minima (C) | Temperatura Maxima (C) | Precipitacao (mm) | Final de Semana | Consumo de cerveja (litros) |
|---|------------|-----------------------|------------------------|------------------------|-------------------|-----------------|-----------------------------|
| 0 | 01/01/2015 | 27,3 | 23,9 | 32,5 | 0 | 0 | 25.461 |
| 1 | 02/01/2015 | 27,02 | 24,5 | 33,5 | 0 | 0 | 28.972 |
| 2 | 03/01/2015 | 24,82 | 22,4 | 29,9 | 0 | 1 | 30.814 |
| 3 | 04/01/2015 | 23,98 | 21,5 | 28,6 | 1,2 | 1 | 29.799 |
| 4 | 05/01/2015 | 23,82 | 21 | 28,3 | 0 | 0 | 28.900 |

rename all the column names

```
In [232... df.columns = ['Date', 'Medium Temp', 'Min Temp', 'Max Temp', 'Precipitation(mm)', 'End of week', 'Beer consumption (liters)']
df.head()
# for renaming a specific column- df.rename(columns = {'old1':'new2', 'old2':new2}, inplace=True)
```

```
Out[232... 
```

| | Date | Medium Temp | Min Temp | Max Temp | Precipitation(mm) | End of week | Beer consumption (liters) |
|---|------------|-------------|----------|----------|-------------------|-------------|---------------------------|
| 0 | 01/01/2015 | 27,3 | 23,9 | 32,5 | 0 | 0 | 25.461 |
| 1 | 02/01/2015 | 27,02 | 24,5 | 33,5 | 0 | 0 | 28.972 |
| 2 | 03/01/2015 | 24,82 | 22,4 | 29,9 | 0 | 1 | 30.814 |
| 3 | 04/01/2015 | 23,98 | 21,5 | 28,6 | 1,2 | 1 | 29.799 |
| 4 | 05/01/2015 | 23,82 | 21 | 28,3 | 0 | 0 | 28.900 |

```
In [233... df.info() # getting the dataframe information like null count and datatype of each column.
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  365 non-null   object
1   Medium Temp           365 non-null   object
2   Min Temp              365 non-null   object
3   Max Temp              365 non-null   object
4   Precipitation(mm)     365 non-null   object
5   End of week           365 non-null   int64
6   Beer consumption (liters) 365 non-null   float64
```

```
dtypes: float64(1), int64(1), object(5)
memory usage: 20.1+ KB
```

Replacing , with . in values of continuous variables of the dataframe

```
In [234... df['Medium Temp'].replace(',', '.', inplace=True, regex=True)
```

```
In [235... df['Min Temp'].replace(',', '.', inplace=True, regex=True)
```

```
In [237... df['Max Temp'].replace(',', '.', inplace=True, regex=True)
```

```
In [238... df['Precipitation(mm)'].replace(',', '.', inplace=True, regex=True)
```

```
In [239... df.head()
```

```
Out[239...
      Date  Medium Temp  Min Temp  Max Temp  Precipitation(mm)  End of week  Beer consumption (liters)
0  01/01/2015         27.3      23.9      32.5                0            0             25.461
1  02/01/2015        27.02      24.5      33.5                0            0             28.972
2  03/01/2015        24.82      22.4      29.9                0            1             30.814
3  04/01/2015        23.98      21.5      28.6                1.2            1             29.799
4  05/01/2015        23.82       21      28.3                0            0             28.900
```

changing the datatypes of continuous variable from object to float

```
In [240... df['Medium Temp'] = df['Medium Temp'].astype('float')
```

```
In [241... df['Min Temp'] = df['Min Temp'].astype('float')
```

```
In [242... df['Max Temp'] = df['Max Temp'].astype('float')
```

```
In [243... df['Precipitation(mm)'] = df['Precipitation(mm)'].astype('float')
```

Splitting the 'Date' variable in separate columns of Day and Month

```
In [244... df['Day'] = pd.to_datetime(df.Date, format="%d/%m/%Y").dt.day
```

```
In [245... df['Month'] = pd.to_datetime(df.Date, format = "%d/%m/%Y").dt.month
```

```
In [246... df.info() # getting the dataframe information after above changes
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   Date                  365 non-null   object
1   Medium Temp           365 non-null   float64
2   Min Temp              365 non-null   float64
3   Max Temp              365 non-null   float64
4   Precipitation(mm)     365 non-null   float64
5   End of week           365 non-null   int64
6   Beer consumption (liters) 365 non-null   float64
7   Day                   365 non-null   int64
8   Month                 365 non-null   int64
dtypes: float64(5), int64(3), object(1)
memory usage: 25.8+ KB
```

Dropping the Date column, since is it splitted in Day and Month

```
In [247... df.drop(["Date"], axis = 1, inplace = True)
```

Describing the data


```
In [248... df.describe() # getting the statistical information from the continuous variables of dataframes
```

```
Out[248...
```


| | Medium Temp | Min Temp | Max Temp | Precipitation(mm) | End of week | Beer consumption (liters) | Day | Month |
|-------|-------------|------------|------------|-------------------|-------------|---------------------------|------------|------------|
| count | 365.000000 | 365.000000 | 365.000000 | 365.000000 | 365.000000 | 365.000000 | 365.000000 | 365.000000 |
| mean | 21.226356 | 17.461370 | 26.611507 | 5.196712 | 0.284932 | 25.401367 | 15.720548 | 6.526027 |
| std | 3.180108 | 2.826185 | 4.317366 | 12.417844 | 0.452001 | 4.399143 | 8.808321 | 3.452584 |
| min | 12.900000 | 10.600000 | 14.500000 | 0.000000 | 0.000000 | 14.343000 | 1.000000 | 1.000000 |
| 25% | 19.020000 | 15.300000 | 23.800000 | 0.000000 | 0.000000 | 22.008000 | 8.000000 | 4.000000 |
| 50% | 21.380000 | 17.900000 | 26.900000 | 0.000000 | 0.000000 | 24.867000 | 16.000000 | 7.000000 |
| 75% | 23.280000 | 19.600000 | 29.400000 | 3.200000 | 1.000000 | 28.631000 | 23.000000 | 10.000000 |
| max | 28.860000 | 24.500000 | 36.500000 | 94.800000 | 1.000000 | 37.937000 | 31.000000 | 12.000000 |

Plotting some of the variables to find the relation with the Beer consumption variable


```
In [249... plt.scatter(df['Medium Temp'], df['Beer consumption (liters)'])  
plt.xlabel('Medium Temperature')  
plt.ylabel('Beer consumption in litres')  
plt.title('Beer consumption based on Medium Temperature');
```



```
In [250... plt.scatter(df['Min Temp'], df['Beer consumption (liters)'])  
plt.xlabel('Minimum Temperature')  
plt.ylabel('Beer consumption in litres')  
plt.title('Beer consumption based on Minimum Temperature');
```



```
In [251... plt.scatter(df['Max Temp'], df['Beer consumption (liters)'])  
plt.xlabel('Maximum Temperature')  
plt.ylabel('Beer consumption in litres')  
plt.title('Beer consumption based on Maximum Temperature');
```



```
In [252... df.corr() # gives the correlation among each continuous variable of the dataframe
```

```
Out[252...
```

| | Medium Temp | Min Temp | Max Temp | Precipitation(mm) | End of week | Beer consumption (liters) | Day | Month |
|---------------------------|-------------|-----------|-----------|-------------------|-------------|---------------------------|-----------|-----------|
| Medium Temp | 1.000000 | 0.862752 | 0.922513 | 0.024416 | -0.050803 | 0.574615 | 0.012382 | -0.103169 |
| Min Temp | 0.862752 | 1.000000 | 0.672929 | 0.098625 | -0.059534 | 0.392509 | -0.011206 | -0.172923 |
| Max Temp | 0.922513 | 0.672929 | 1.000000 | -0.049305 | -0.040258 | 0.642672 | 0.035079 | -0.074866 |
| Precipitation(mm) | 0.024416 | 0.098625 | -0.049305 | 1.000000 | 0.001587 | -0.193784 | -0.003414 | 0.007089 |
| End of week | -0.050803 | -0.059534 | -0.040258 | 0.001587 | 1.000000 | 0.505981 | 0.006254 | -0.006526 |
| Beer consumption (liters) | 0.574615 | 0.392509 | 0.642672 | -0.193784 | 0.505981 | 1.000000 | 0.025969 | 0.039908 |
| Day | 0.012382 | -0.011206 | 0.035079 | -0.003414 | 0.006254 | 0.025969 | 1.000000 | 0.011893 |
| Month | -0.103169 | -0.172923 | -0.074866 | 0.007089 | -0.006526 | 0.039908 | 0.011893 | 1.000000 |

Starting creating the Model

splitting the data into predictor(x) and target (y)

```
In [253... x = df.drop('Beer consumption (liters)',axis=1)
y = df['Beer consumption (liters)']
```

```
In [254... from sklearn.model_selection import train_test_split
```

creating the train and test data of the predictor and target

```
In [255... x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

creating the linear regression model

```
In [256... lrm = LinearRegression()
```

```
In [257... lrm.fit(x_train, y_train)
```

```
Out[257... LinearRegression()
```

```
In [258... lrm.score(x_train, y_train) #training score of the created model
```

```
Out[258... 0.7222896489940203
```

```
In [259... lrm.score(x_test, y_test) #test score of the created model
```

Out[259... 0.7406013522605527

Trying feature selection to select the best variables for prediction and try increasing the model accuracy

```
In [260... from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2, f_regression
```

```
In [261... bestfeatures = SelectKBest(score_func=f_regression, k='all')
```

```
In [262... fit = bestfeatures.fit(x,y)
```

```
In [263... fit.scores_
```

Out[263... array([1.78938298e+02, 6.61100241e+01, 2.55428041e+02, 1.41633735e+01,
1.24913906e+02, 2.44963789e-01, 5.79047775e-01])

```
In [264... x.columns
```

Out[264... Index(['Medium Temp', 'Min Temp', 'Max Temp', 'Precipitation(mm)',
'End of week', 'Day', 'Month'],
dtype='object')

```
In [265... dfscores = pd.DataFrame(fit.scores_)
dfcolumns = pd.DataFrame(x.columns)
```

```
In [266... featureScores = pd.concat([dfcolumns,dfscores],axis=1)
featureScores
```

Out[266...

| | | 0 | 0 |
|---|-------------------|------------|---|
| 0 | Medium Temp | 178.938298 | |
| 1 | Min Temp | 66.110024 | |
| 2 | Max Temp | 255.428041 | |
| 3 | Precipitation(mm) | 14.163373 | |
| 4 | End of week | 124.913906 | |
| 5 | Day | 0.244964 | |
| 6 | Month | 0.579048 | |

```
In [267... featureScores.columns = ['effect_on_consumption','Score']
featureScores
```

Out[267...

| | effect_on_consumption | Score |
|---|-----------------------|------------|
| 0 | Medium Temp | 178.938298 |
| 1 | Min Temp | 66.110024 |
| 2 | Max Temp | 255.428041 |
| 3 | Precipitation(mm) | 14.163373 |
| 4 | End of week | 124.913906 |
| 5 | Day | 0.244964 |
| 6 | Month | 0.579048 |

```
In [268... featureScores.nlargest(3,'Score') #getting these best three variable among 7 after feature selection
```

Out[268...

| | effect_on_consumption | Score |
|---|-----------------------|------------|
| 2 | Max Temp | 255.428041 |

| | | |
|---|-------------|------------|
| 0 | Medium Temp | 178.938298 |
| 4 | End of week | 124.913906 |

```
In [269... x1 = df[['Max Temp','Medium Temp','End of week']] #predictor suggested by feature selection
y1 = df['Beer consumption (liters)']
```

```
In [270... x1_train, x1_test, y1_train, y1_test = train_test_split(x1,y1,test_size=0.2,random_state=42)
```

```
In [271... lrm1 = LinearRegression()
```

```
In [272... lrm1.fit(x1_train, y1_train)
```

```
Out[272... LinearRegression()
```

```
In [273... lrm1.score(x1_train, y1_train)
```

```
Out[273... 0.6954059612625987
```

```
In [274... lrm1.score(x1_test, y1_test) # accuracy didn't improve using the suggested 3 variables after feature selection
```

```
Out[274... 0.6886564844629807
```

```
In [275... featureScores.nlargest(5,'Score') #getting top 5 features and will try model creation using these features to get
```

```
Out[275...
```

| | effect_on_consumption | Score |
|---|-----------------------|------------|
| 2 | Max Temp | 255.428041 |
| 0 | Medium Temp | 178.938298 |
| 4 | End of week | 124.913906 |
| 1 | Min Temp | 66.110024 |
| 3 | Precipitation(mm) | 14.163373 |

```
In [276... x1 = df[['Max Temp','Medium Temp','End of week','Min Temp','Precipitation(mm)']] #predictor suggested by feature
y1 = df['Beer consumption (liters)']
```

```
In [277... x1_train, x1_test, y1_train, y1_test = train_test_split(x1,y1,test_size=0.2,random_state=42)
```

```
In [278... lrm1 = LinearRegression()
lrm1.fit(x1_train, y1_train)
```

```
Out[278... LinearRegression()
```

```
In [279... lrm1.score(x1_train, y1_train)
```

```
Out[279... 0.7104068706547428
```

```
In [280... lrm1.score(x1_test, y1_test) # Accuracy score increased by 6 percent after using 5 variables (feature select) in t
```

```
Out[280... 0.7427938181737366
```

Trvino Standard Scaling on train and test data of all predictors in order to trv getting more accuracv

Trying standard scaling on train and test data of all predictors in order to try getting more accuracy,

```
In [281... from sklearn.preprocessing import StandardScaler
```

```
In [282... s = StandardScaler()
```

```
In [283... x_train = s.fit_transform(x_train)
x_test = s.fit_transform(x_test)
```

```
In [284... col = x.columns
```

```
In [285... x_train = pd.DataFrame(x_train, columns=[col])
x_test = pd.DataFrame(x_test, columns=[col])
x_train.head()
```

```
Out[285...
   Medium Temp  Min Temp  Max Temp  Precipitation(mm)  End of week   Day   Month
0    -0.577526  -1.569856   0.031792        -0.425154   -0.593171  1.725816  0.125584
1    -0.329620   0.237186  -0.590171         3.512292   -0.593171 -0.992053  1.616253
2    -2.040166  -1.569856  -2.386953        -0.425154   -0.593171  0.706615  0.125584
3     0.252957   0.445691  -0.198564         1.568281   -0.593171  0.140392 -1.066952
4     0.581431   0.341439   0.768934        -0.425154   -0.593171  1.159593 -1.066952
```

```
In [286... lrm2 = LinearRegression()
```

```
In [287... lrm2.fit(x_train, y_train)
```

```
Out[287... LinearRegression()
```

```
In [288... lrm2.score(x_train, y_train)
```

```
Out[288... 0.7222896489940203
```

```
In [289... lrm2.score(x_test, y_test) # scaling didn't worked in getting the more accuracy, on the contrary the accuracy dec
```

```
Out[289... 0.6777625864500789
```

Trying the different regressors

```
In [290... X = df.drop('Beer consumption (liters)',axis=1)
Y = df['Beer consumption (liters)']
```

```
In [291... X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

```
In [292... from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import AdaBoostRegressor
from sklearn.svm import SVR
```

```
In [293... model = [DecisionTreeRegressor,SVR,RandomForestRegressor,KNeighborsRegressor,AdaBoostRegressor]

for mod in model:
    reg = mod()
    reg = reg.fit(X_train,Y_train)
    print(mod , 'accuracy',reg.score(X_test,Y_test))
```

```

<class 'sklearn.tree._classes.DecisionTreeRegressor'> accuracy 0.47627933938585965
<class 'sklearn.svm._classes.SVR'> accuracy 0.359657180864009
<class 'sklearn.ensemble._forest.RandomForestRegressor'> accuracy 0.6931273024263127
<class 'sklearn.neighbors._regression.KNeighborsRegressor'> accuracy 0.3153496691022142
<class 'sklearn.ensemble._weight_boosting.AdaBoostRegressor'> accuracy 0.7031416451200945

```

```

In [294... # on scaled data
model_s = [DecisionTreeRegressor,SVR,RandomForestRegressor,KNeighborsRegressor,AdaBoostRegressor]

for mod in model_s:
    reg = mod()
    reg = reg.fit(x_train,y_train)
    print(mod , 'accuracy',reg.score(x_test,y_test))

<class 'sklearn.tree._classes.DecisionTreeRegressor'> accuracy 0.23763566840738326
<class 'sklearn.svm._classes.SVR'> accuracy 0.6621135292581894
<class 'sklearn.ensemble._forest.RandomForestRegressor'> accuracy 0.6727914207125629
<class 'sklearn.neighbors._regression.KNeighborsRegressor'> accuracy 0.7161239938784414
<class 'sklearn.ensemble._weight_boosting.AdaBoostRegressor'> accuracy 0.6988167270799108

```

None of the above regressor has better score than the linear regression model created earlier.

creating the linear regression model with combinations of variables in the predictor to get the best accuracy

```

In [297... x3 = df.drop(['Beer consumption (liters)','Day','Month'],axis=1)
y3 = df['Beer consumption (liters)']

```

```

In [298... x3_train, x3_test, y3_train, y3_test = train_test_split(x3, y3, test_size=0.2, random_state=42)

```

```

In [299... lrm3 = LinearRegression()

```

```

In [300... lrm3.fit(x3_train, y3_train)

```

```

Out[300... LinearRegression()

```

```

In [301... lrm3.score(x3_train, y3_train)

```

```

Out[301... 0.7104068706547428

```

```

In [302... lrm3.score(x3_test, y3_test)

```

```

Out[302... 0.7427938181737361

```

```

In [325... x4 = df[['Medium Temp','Max Temp','Min Temp','End of week']]
y4 = df['Beer consumption (liters)']

```

```

In [326... x4_train, x4_test, y4_train, y4_test = train_test_split(x4, y4, test_size=0.2, random_state=42)

```

```

In [327... lrm4 = LinearRegression()

```

```

In [328... lrm4.fit(x4_train, y4_train)

```

```

Out[328... LinearRegression()

```

```

In [329... lrm4.score(x4_train, y4_train)

```

```

Out[329... 0.6961296970163136

```



```
In [330... lrm4.score(x4_test, y4_test)
```

```
Out[330... 0.6861100648174152
```

```
In [309... x5 = df[['Max Temp', 'Medium Temp', 'End of week']]  
y5 = df['Beer consumption (liters)']
```

```
In [310... x5_train, x5_test, y5_train, y5_test = train_test_split(x5, y5, test_size=0.2, random_state=42)
```

```
In [311... lrm5 = LinearRegression()  
lrm5.fit(x5_train, y5_train)
```

```
Out[311... LinearRegression()
```

```
In [312... lrm5.score(x5_train, y5_train)
```

```
Out[312... 0.6954059612625987
```

```
In [313... lrm5.score(x5_test, y5_test)
```

```
Out[313... 0.6886564844629807
```

After trying multiple combinations for predictor variables , got the best score for below predictor variables combination.

```
In [314... x6 = df[['Medium Temp', 'End of week', 'Max Temp', 'Precipitation(mm)']] # predictor  
y6 = df['Beer consumption (liters)'] # target
```

```
In [315... x6_train, x6_test, y6_train, y6_test = train_test_split(x6, y6, test_size=0.2, random_state=42)
```

```
In [316... lrm6 = LinearRegression()  
lrm6.fit(x6_train, y6_train)
```

```
Out[316... LinearRegression()
```

model training accuracy

```
In [317... lrm6.score(x6_train, y6_train)
```

```
Out[317... 0.7101409516827788
```

model testing accuracy

```
In [318... lrm6.score(x6_test, y6_test)
```

```
Out[318... 0.7443500987830509
```

```
In [320... y6_prediction = lrm6.predict(x6_test) # prediction
```

mean squared error

mean squared error

In [321...

```
print('mean squared error is: ', mean_squared_error(y6_test, y6_prediction))
```

mean squared error is: 5.66496605013936

mean absolute error

In [322...

```
print('mean absolute error is: ', mean_absolute_error(y6_test, y6_prediction))
```

mean absolute error is: 2.05538301484382

root mean squared error

In [323...

```
print('root mean squared error is: ', np.sqrt(mean_squared_error(y6_test, y6_prediction)))
```

root mean squared error is: 2.3801189151257462

r2 score

In [324...

```
print('r2 score is: ', r2_score(y6_test, y6_prediction))
```

r2 score is: 0.7443500987830509